

Development to Deployment

- Some bugs only appear under stress
- Prod environment is not the same as the dev env
- Deployment exposes your app in ways your app was not meant to be used
 - Evil forces
 - Idiots

Early Deployment Strategies

- Need a Virtual Private Server (VPS)
 - Might be in cloud eg: AWS
- Install and configure OS and all utilities needed for the app
- Fix almost-weekly security vulnerabilities
- Tune all moving parts to get most bang for buck
- Figure out how to automate horizontal scaling

Platform as a Service (PaaS)

- Easy tiers of horizontal scaling
- Component-level performance tuning
- Infrastructure-level security
- App devs still have to worry about usage and efficiency of app

Performance Metrics

- Availability / uptime \triangleq percent of time site is up and accessible
- Responsiveness \triangleq how long after a click does user get response
- Scalability \triangleq as users increase, can you maintain responsiveness without increasing cost/user

Availability

- Measured as a percent of uptime over a given window
- Defined in terms of 9's
 - ‘five nines’ \triangleq 99.999% uptime
- Need to consider usage and implementation metrics as well
 - Server vs cluster availability
 - Availability at load

Responsiveness

- <100ms is ‘instantaneous’
- 7sec is abandonment
- Want to consider what most users see; cannot model on a standard distribution
- Scalability
 - As number of users increase, response time to each user stays the same
 - Cost to serve each user stays the same
- Service Level Objective (SLO) \triangleq percent of users get acceptable performance
 - Specify percentile, target response time, and time window
- Service Level Agreement (SLA) \triangleq an SLO to which provider is contractually obligated
- Application Performance Index (apdex) \triangleq defines satisfactory performance

- Given a threshold latency T for user satisfaction
- Satisfactory requests take $t \leq T$
- Tolerable requests take $T \leq t \leq 4T$
- $\text{Apdex} := \frac{|\text{satisfactory}| + \frac{1}{2}|\text{tolerable}|}{|\text{requests}|}$
- $\mathbb{E}[\text{apdex}|\text{good}] \in [0.85, 0.93]$
- Increase responsiveness
 - Small sites: overprovision presentation and logic tier
 - Large site: worry...
- Response time $:= t(\text{render done}) - t(\text{request sent})|_{t(x) \triangleq \text{time at event } x}$

Security Metrics

- Privacy \triangleq is data access limited to the appropriate users
- Authentication \triangleq can we trust that user is who s/he claims to be
- Data integrity \triangleq is user's sensitive data tamper-evident

Three-Tier Architecture

- Type of shared nothing architecture
1. Presentation tier
 2. Logic (application) tier
 3. Persistence tier

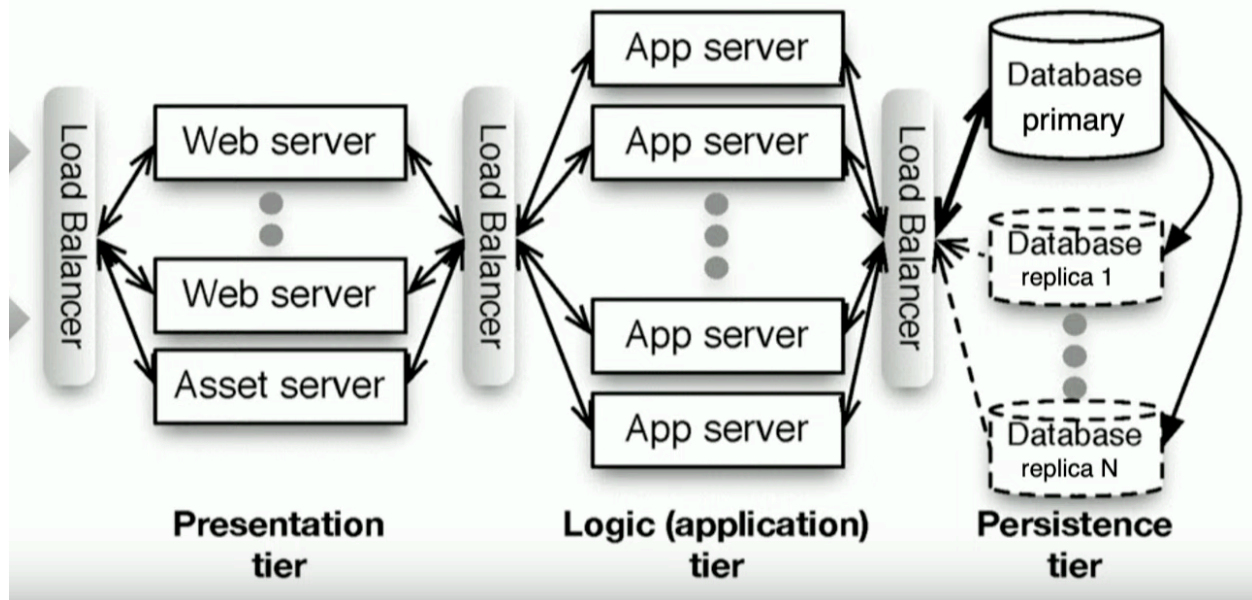


Figure 1: 475

Dynamic Content Generation

- Run a program to generate each 'page'
- Originally, web used template with embedded code 'snippets'
- Eventually, code moved out of the web server

Database Scaling

Sharding

- Partition across independent ‘shards’
- Scales great
- Bad when operations touch > 1 table
- Eg: good for user profile info

Replication

- Replicate all data everywhere
- Good for multi-table queries
- Hard to scale b/c writes must propagate to all copies \implies temporary inconsistency in data values
- Example: social media ‘wall’ posts

Parts of PaaS App Controlled by Dev

- App (controller and model)
- Database access
 - Avoid abusive queries
 - Use indices wisely
- View rendering
 - Use CSS selectors wisely (`id` and `class` is fast, but DOM traversal isn’t)
 - Cache ‘expensive’ fragments / partials in views
- Embedded assets (images, etc.) \triangleq serve from dedicated static assets servers
- JavaScript \triangleq load popular libraries from asset servers with generous expiration

CICD

Continuous Integration (CI)

- Integration testing the app beyond what each developer does
- Pre-release code checkin triggers CI
- Since frequent checkins, CI always running
- Common strategy: integrate with GitHub
- Good for:
 - Diagnosing differences between dev and prod envs
 - Cross-browser or cross-version testing
 - Testing SOA integration when remote services don’t act as planned
 - Hardening / protecting against attacks
 - Stress testing / longevity testing of new features / code paths

Continuous Deployment (CD)

- Lots of deploys, very often
- Rational risk \triangleq number of engineer-hours invested in product since last deploy
- Deployment should be a non-event; something that happens all of the time
 - Automate when possible
- Can configure to be deployed on push
 - Can be auto-integrated with CI
- Good for smaller features and non-customer facing changes
- Bad for (use releases for) big changes in how the user interacts with the app
 - Can tag commits to define releases with `git tag`
 - Have to push tags with `git push --tags`