

Uncertainty and Randomness (Expectimax)

- Good for explicit randomness, unpredictable opponents, or actions that can fail
- Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcomes
- Expectimax \triangleq compute the average scores under optimal play
 - Max nodes as in minimax search
 - Chance nodes are like min nodes, but the outcome is uncertain
 - Calculate their expected utilities \triangleq weighted average (expectation) of children

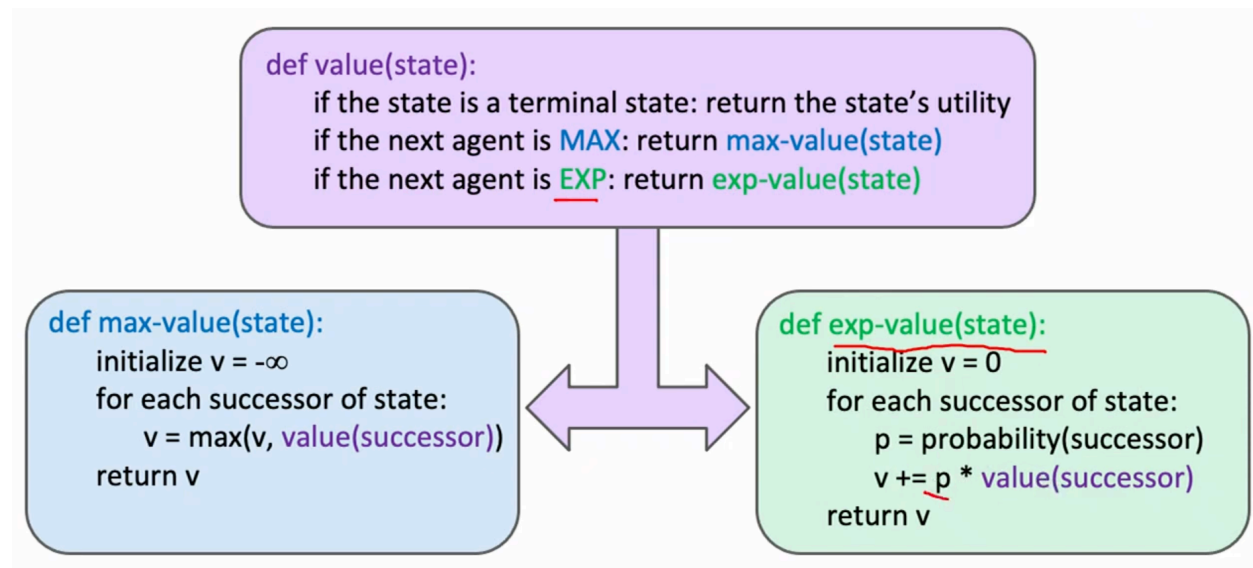


Figure 1: Screenshot_2023-09-20_at_6.51.31_PM.png

-
- Cannot prune
- Depth-limitation works
- Expected value of a function of a random variable \triangleq the average weighted by the probability distribution over outcomes
- Can be challenging to get the optimal weights for an expectimax model

Mixed Layer Types

- Expectiminimax
 - Environment is an extra ‘random agent’ player that moves after each min/max agent
 - Each node computes the appropriate combination of its children

Multi-Agent Utilities

- Minimax
 - Have layers for each player
 - Have to propagate each value for each player up the tree
 - Terminals have utility tuples
 - Each player maximizes its own value

Difficulties with Search

- Even with alpha-beta pruning and limited depth, large b is an issue ($O(b^{\frac{m}{2}})$)
- Limiting depth requires us to design good evaluation functions
 - Not a general-purpose method: need to design new evaluation function for each new problem
 - Bad evaluation function may lead to inefficient or biased solutions
 - The trend in AI is to prefer general method and less human tweaking

Monte Carlo Tree Search (MCTS)

- Evaluation by rollouts: estimate value of a state by playing many games from this state by taking random actions (or some other fast policy) and count wins and losses
 - Allocate more rollouts to promising or uncertain nodes
- Selective search: explore parts of the tree that will help improve the decision at the root, regardless of depth

Upper Confidence Bound (UCB) Heuristic

$$UCB1(n) = \frac{U(n)}{N(n)} + C \times \sqrt{\frac{\log(\text{parent}(n))}{N(n)}} \quad \left| \begin{array}{l} C \hat{=} \text{parameter we chose to trade off between two terms} \\ N(n) \hat{=} \text{number of rollouts from node } n \\ U(n) \hat{=} \text{total utility of rollouts for player of parent}(n) \end{array} \right.$$

- The total utility of rollouts $\hat{=}$ the number of wins - Have to keep track of both N and U for each node - Combines ‘promising’ and ‘uncertain’

Monte Carlo Tree Search (MCTS) Algorithm

- Repeat until end of time:
 1. Selection: recursively apply UCB to choose a path down and leaf node n
 2. Expansion: add a new child c to n
 3. Simulation: run a rollout from c
 4. Back-propagation: update U and N counts from c back up to the root
- Choose the action leading to the child with highest N
- Most common tool for solving hard search problems
- Time complexity independent of b and m
- No need to design evaluation functions (general-purpose and easy to use)
- Solution quality depends on number of rollouts N
 - Theorem: as $N \rightarrow \infty$, MCTS selects the minimax move
- Maps well to parallel hardware