

## Reinforcement Learning Contd.

### Passive RL

- How to learn from already given experiences
- Given:
  - Set of states  $s \in S$
  - Set of actions (per state)  $A$
- Looking for  $\pi(s)$
- Don't know
  - A model  $T(s, a, s')$
  - A reward function  $R(s, a, s')$
- Big idea  $\triangleq$  compute all averages over  $T$  using sample outcomes

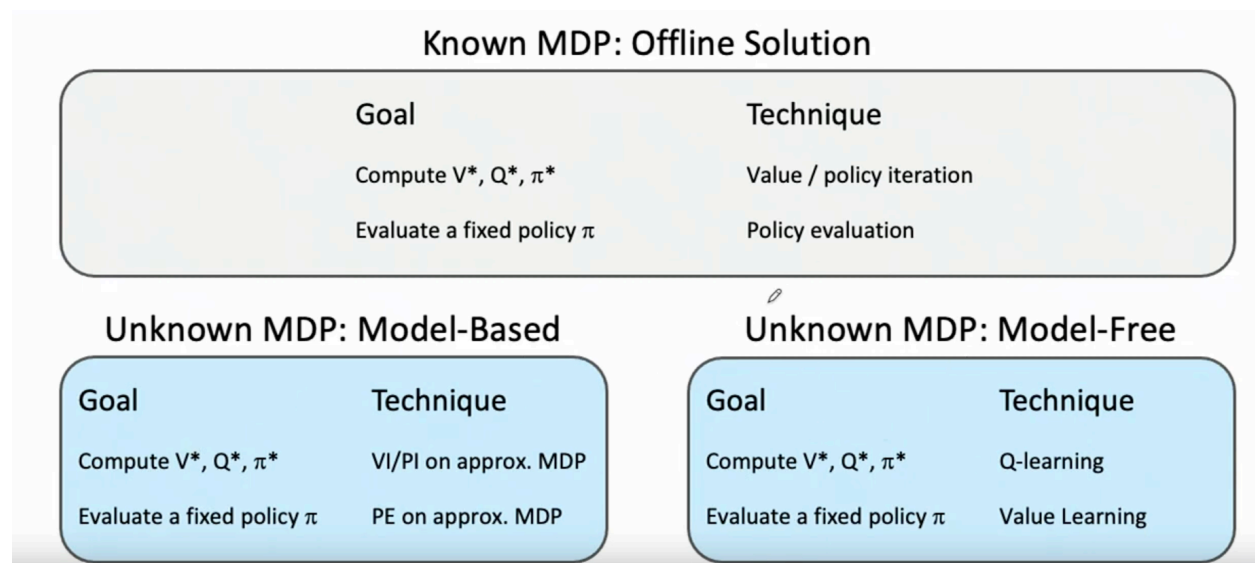


Figure 1: Screenshot\_2023-10-02\_at\_5.18.04\_PM.png

•

### Model-Free Learning

- Temporal difference learning
- Receives stream of experiences from the world eg:  $(s, a, r, s', a', r', s'', \dots)$
- Update estimates each transition
- Over time, updates will mimic Bellman updates
- Q-iteration: do Q-value updates to each Q-state:

$$Q_{k+1}(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \max_{a'} Q_k(s', a')]$$

- $Q_0(s, a) = 0$
- Can't compute this update without  $T$  and  $R$
- Q-learning: instead, compute average as we go
  - Receive a sample transition  $(s, a, r, s')$

- Use initial approximation  $Q(s, a) \approx r + \gamma \max_{a'} Q(s', a')$  to compute:

$$Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a')]$$

- $\pi^*(s) = \arg \max_a (Q(s, a))$
- AKA off-policy learning
- Have to explore enough to be good
- Have to eventually make the learning rate small enough (but not decrease it too quickly)
- Type of active RL

## Active RL

- Need to consider exploration vs exploitation

### Choosing How to Explore

- Simplest: random actions ( $\epsilon$ -greedy)
  - Every time step, flip a coin
  - With (small) probability  $\epsilon$ , act randomly
  - With (large) probability  $1 - \epsilon$ , act on current policy
  - $\epsilon$  high during early learning stage and low at end
- Better idea: explore areas whose badness is not (yet) established; eventually stop exploring
  - Use an exploration function that takes a value estimate  $u$  and a visit count  $n$  and returns an optimistic utility, eg:  $f(u, n) = u + \frac{k}{n} | k \triangleq \text{preset constant}$
  - Then use

$$Q(s, a) \leftarrow_{\alpha} \alpha R(s, a, s') + \gamma \max_{a'} f(Q(s', a'), N(s', a'))$$

\* Note:  $x \leftarrow_{\alpha} \implies x \leftarrow (1 - \alpha)x + \alpha v$

### Regret

- Even if you learn the optimal policy, you still make mistakes along the way
- Regret  $\triangleq$  a measure of total mistake cost
  - Difference between your (expected) rewards, including youthful suboptimality, and optimal (expected) rewards
- eg: random exploration and exploration functions both end up optimal, but random exploration has higher regret

## Dealing with Large State Spaces

- Use approximate Q-learning
- Basic Q-Learning keeps a table of all q-values
  - This is not possible in reality; too much state
- Idea: Learn about some small number of training states from experience
  - Generalize that experience to new, similar situations
- Solution: linear value functions: describe state using a vector of features (properties)  $(f_1, f_2, f_3, \dots, f_n)$ 
  - Features are functions from states to real numbers (often 0/1) that capture important properties of the state
    - \* eg: distance to closest ghost, is pacman in a tunnel, etc.
    - \* Can also describe a q-state  $(s, a)$  with features
- Using a feature representation, we can write a Q-function (or value function) for any state using a few weights  $(w_1, w_2, \dots, w_n)$ :

$$V(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Advantage: our experience is summed up in a few powerful numbers  $\{w_1, w_2, \dots, w_n\}$
- Disadvantage: states may share features but actually be very different in value
- Q-learning with linear Q-functions:

$$\begin{aligned}
 \text{transition} &= (s, a, r, s') \\
 \text{difference} &= [r + \gamma \max_{a'} Q(s', a')] - Q(s, a) \\
 Q(s, a) &\leftarrow Q(s, a) + \alpha[\text{difference}] \\
 w_i &\leftarrow w_i + \alpha[\text{difference}]f_i(s, a)
 \end{aligned}$$