

Machine Learning

- ML := how to acquire a model from data / experience
 - Learning parameters (eg: probabilities)
 - Learning structure (eg: BN graphs)
 - Learning hidden concepts (eg: clustering)
- Given training data \triangleq example data points and their actual values
- Goal \triangleq build model based on training data that outputs results in a similar fashion

Types of Learning Problems


- Supervised learning \triangleq correct answers for each training instance
 - Classification \triangleq learning predictor with discrete outputs
 - Regression \triangleq learning predictor with real-valued outputs
- Reinforcement learning \triangleq reward sequence, no correct answers
- Unsupervised learning \triangleq ‘just make sense of the data’

Classification

- Dataset := each data point, x , is associated with some label (aka class), y
- Goal \triangleq given inputs x , write alg to predict labels y
- Workflow:
 - x given
 - Extract features from input \triangleq attributes of the input that characterize x and hopefully help with classification
 - Run some machine learning alg on the features
 - Output a predicted label y

Model-based Approach

- Build a model (eg: BN) where both the labels and features are random variables
- Instantiate any observed features
- Query for the distribution of the label conditioned on the features
- Challenges: how to structure BN / how to learn parameters



- Y := the label
- $\{F_1, F_2, \dots, F_n\}$:= the n features
- Prior := $\mathbb{P}(Y)$
- Generate a table per feature $\mathbb{P}(F_i|Y)$
 - Collectively called ‘parameters’; denoted as θ
- Training \triangleq use the training dataset to estimate the probability tables:
 1. Estimate $\mathbb{P}(Y) \triangleq$ how often does each label occur
 2. Estimate $\mathbb{P}(F_i|Y) \triangleq$ how does the label affect the feature

- Classification \triangleq instantiate all features given the input features as evidence
 1. Derive features from sample
 2. Query for $\mathbb{P}(Y|F_1, F_2, \dots, F_n) \triangleq$ probability of label, given all the input features
 - Use an inference alg (eg: variable elimination) to compute this

Naive Bayes

1. Extract features
 2. Generate BN
 3. Get joint probability of label and evidence for each label $\mathbb{P}(Y, f_1, \dots, f_n)$
 4. Sum to get probability of evidence
 5. Normalize
 6. Return the y with maximal probability
 - Can restrict to a confidence bound
- Good because returns exact confidence bound for every possible y
 - Generally $\mathbb{P}(Y, F_1, \dots, F_n) = \mathbb{P}(Y) \prod_i \mathbb{P}(F_i|Y)$
 - Works because features are assumed to be independent
 - $n \times |F| \times |Y|$ parameters (linear in n)

Bag-of-words Model

- Each feature is identically distributed
- All features share the same conditional probs $\mathbb{P}(W|Y)$

Parameter Estimation

- Estimating the distribution of a random variable
- Elicitation \triangleq ask a human
- Empirically \triangleq use training data

Maximum Likelihood

- Chose the θ value that maximizes the probability of the observation
 - Find θ that maximizes $\mathbb{P}(\text{observation}|\theta)$
- eg:
 - Given $\mathbb{P}(\text{red}|\theta) = \theta$
 - $\mathbb{P}(2 \text{ red and } 1 \text{ blue}|\theta \text{ of beans are red}) = \mathbb{P}(\text{red}|\theta)\mathbb{P}(\text{red}|\theta)\mathbb{P}(\text{blue}|\theta) = \theta^2(1 - \theta)$
 - Want to compute $\theta^* = \arg \max_{\theta} \theta^2(1 - \theta)$
 - * Find θ such that $\frac{d}{d\theta} \theta^2(1 - \theta) = 0$
- Note: maximizing the likelihood is the same as maximizing the log-likelihood
 - $\arg \max_{\theta} f(\theta) = \arg \max_{\theta} \ln f(\theta)$
 - Makes derivatives much easier because it turns products into sums
- For naive BNs:
 - $\mathbb{P}(y) = \frac{\# \text{ occurrences of class } y}{\text{total } \# \text{ of observations}}$
 - $\mathbb{P}(f|y) = \frac{\# \text{ of occurrences of feature } f \text{ and class } y}{\text{total } \# \text{ of occurrences of class } y}$